

Computer Organization and Architecture

Module 5 (Part 2)

Design of Memory Subsystems

Prof. Indranil Sengupta

Dr. Sarani Bhattacharya


Department of Computer Science and Engineering

IIT Kharagpur

Memory Interfacing and Addressing


Memory Interfacing

- Basic problem:
 - Interfacing one or more memory modules to the processor.
 - We assume a single level memory at present (i.e. no cache memory).
- Questions to be answered:
 - How the processor address and data lines are connected to memory modules?
 - How are the addresses decoded?
 - How are the memory addresses distributed among the memory modules?
 - How to speed up data transfer rate between processor and memory?



The processor's view of memory

- Typical interface of a memory module.
- Real chip may contain more signal lines (e.g. DRAM).



A Note About the Memory Interface Signals


- The data signals of a memory module (RAM) are typically bidirectional.
 - Some memory chips may have separate data in and data out lines.
- For memory ***READ*** operation:
 - Address of memory location is applied to *address lines*.
 - ***RD/WR'*** control signal is set to 1, and ***CS'*** is set to 0.
 - Data is read out through the *data lines* after memory access time delay.
- For memory ***WRITE*** operation:
 - Address of memory location is applied to *address lines*, and the data to be written to *data lines*.
 - ***RD/WR'*** control signal is set to 0, and ***CS'*** is set to 0.

- Why is CS' signal required?

- To handle multiple memory modules interfacing problem.
- We typically select only one out of several memory modules at a time.

- What happens when $CS' = 1$?


- When a memory module is *not selected*, the data lines are set to the *high impedance state* (i.e. electrically disconnected).
- An example scenario is shown.



An Example Memory Interfacing Problem

- Consider a MIPS32 like processor with a 32-bit address.
 - Maximum memory that can be connected is $2^{32} = 4$ Gbytes.
 - Assume that the processor data lines are 8 bits.
- Assume that memory chips (RAM) are available with *size 1 Gbyte*.
 - 30 address lines and 8 data lines.
 - Low-order 30 address lines ($A_{29}-A_0$) are connected to the memory modules.
- We want to interface *4 such chips* to the processor.
 - Total memory of 4 Gbytes.

PROCESSOR



- High order address lines (A_{31} and A_{30}) select one of the memory modules.
- **When is M0 selected?**
 - Address is: **0 0** xxxxxxxx xxxx xxxx xxxx xxxx xxxx xxxx
 - Range of addresses is: 0x**00000000** to 0x**3FFFFFFF**
- **When is M1 selected?**
 - Address is: **0 1** xxxxxxxx xxxx xxxx xxxx xxxx xxxx xxxx
 - Range of addresses is: 0x**40000000** to 0x**7FFFFFFF**
- **When is M2 selected?**
 - Address is: **1 0** xxxxxxxx xxxx xxxx xxxx xxxx xxxx xxxx
 - Range of addresses is: 0x**80000000** to 0x**BFFFFFFF**
- **When is M3 selected?**
 - Address is: **1 1** xxxxxxxx xxxx xxxx xxxx xxxx xxxx xxxx
 - Range of addresses is: 0x**C0000000** to 0x**FFFFFFFF**

- **An observation:**
 - Consecutive block of bytes are mapped to the same memory module.
 - For MIPS32, we have to access 32 bits (4 bytes) of data in parallel, which requires four sequential memory accesses here.
 - We shall look at an alternate memory organization later that would make this possible.
 - Called *memory interleaving*.

Another Example

- Use 1K x 4 memory modules to build a 8K x 8 memory system.

Improved Memory Interface for MIPS32


- We make small changes in the organization so that 32-bits of data can be fetched in a single memory access cycle.
 - Exploit the concept of *memory interleaving*.
 - Consecutive bytes are mapped to different memory modules.
- The main changes:
 - High order 30 address lines (A_{31} - A_2) are connected to memory modules.
 - Low order two address lines (A_1 and A_0) are used to select one of the modules.

- How are the addresses mapped to memory modules?
 - **Module M0**: 0, 4, 8, 12, 16, 20, 24, ...
 - **Module M1**: 1, 5, 9, 13, 17, 21, 25, ...
 - **Module M2**: 2, 6, 10, 14, 18, 22, 26, ...
 - **Module M3**: 3, 7, 11, 15, 19, 23, 27, ...
- Memory addresses are *interleaved* across memory modules.
- What we can gain from this mapping?
 - Consecutive addresses are mapped to consecutive modules.
 - Possible to access four consecutive words in the same cycle, if all four modules are enabled simultaneously.

- Motivation for word alignment in MIPS32 data words.
 - 32-bit words start from a memory address that is divisible by 4.
 - Corresponding byte addresses are (0, 1, 2, 3), (4, 5, 6, 7), (8, 9, 10, 11), (12, 13, 14, 15), etc.
 - Possible to transfer all the four bytes of a word in a *single memory cycle*.
 - What happens if a word is not aligned?
 - Say: (1, 2, 3, 4) or (2, 3, 4, 5) or (3, 4, 5, 6).
 - Two of the bytes will be mapped to the same memory module.
 - Hence the word cannot be transferred in a single memory cycle.


2 memory cycles required

MIPS32 PROCESSOR



Still one module is selected at a time
:: 8 bits data transfer per cycle.

MIPS32 PROCESSOR



Enable all the four modules together.
32-bit parallel data transfer.

Memory Latency and Bandwidth

- **Memory Latency:**

- The delay from the issue of a memory read request to the first byte of data becoming available.

- **Memory Bandwidth:**

- The maximum number of bytes that can be transferred between the processor and the memory system per unit time.

- **Example 1:**

Consider a memory system that takes 20 ns to service the access of a single 32-bit word.

Latency L = 20 ns per 32-bit word.

Bandwidth BW = $32 / (20 \times 10^{-9}) = 200$ Mbytes per second.

- **Example 2:**

The memory system is modified to accept a new (still 20ns) request for a 32-bit word every 5 ns by overlapping requests.

Latency L = 20 ns per 32-bit word (*no change*).

Bandwidth BW = $32 / (5 \times 10^{-9}) = 800$ Mbits per second.

Memory Hierarchy Design

Introduction

- Programmers want unlimited amount of memory with very low latency.
- Fast memory technology is more expensive per bit than slower memory.
 - SRAM is more expensive than DRAM, DRAM is more expensive than disk.
- Possible solution?
 - Organize the memory system in several levels, called *memory hierarchy*.
 - Exploit temporal and spatial locality on computer programs.
 - Try to keep the commonly accessed segments of program / data in the faster memories.
 - Results in faster access times on the average.

Quick Review of Memory Technology

- **Static RAM:**

- Very fast but expensive memory technology (requires 6 transistors / bit).
- Packing density is limited.

- **Dynamic RAM:**

- Significantly slower than SRAM, but much less expensive (1 transistor / bit).
- Requires periodic refreshing.

- **Flash memory:**

- Non-volatile memory technology that uses floating-gate MOS transistors.
- Slower than DRAM, but higher packing density, and lower cost per bit.


- **Magnetic disk:**

- Provides large amount of storage, with very low cost per bit.
- Much slower than DRAM, and also flash memory.
- Requires mechanical moving parts, and uses magnetic recording technology.

Memory Hierarchy

- The memory system is organized in several levels, using progressively faster technologies as we move towards the processor.
 - The entire addressable memory space is available in the largest (but slowest) memory (typically, magnetic disk or flash storage).
 - We incrementally add smaller (but faster) memories, each containing a *subset* of the data stored in the memory below it.
 - We proceed in steps towards the processor.

- Typical hierarchy (starting with closest to the processor):
 1. Processor registers
 2. Level-1 cache (typically divided into separate instruction and data cache)
 3. Level-2 cache
 4. Level-3 cache
 5. Main memory
 6. Secondary memory (magnetic disk / solid-state drive)
- As we move away from the processor:
 - a) Size increases
 - b) Cost decreases
 - c) Speed decreases




A Comparison

Level	Typical Access Time	Typical Capacity	Other Features
Register	300-500 ps	500-1000 B	On-chip
Level-1 cache	1-2 ns	16-64 KB	On-chip
Level-2 cache	5-20 ns	256 KB – 2 MB	On-chip
Level-3 cache	20-50 ns	1-32 MB	On or off chip
Main memory	50-100 ns	1-16 GB	
Magnetic disk	5-50 ms	100 GB – 16 TB	

Processor-Memory Performance Gap

- Processor is much faster than main memory.
 - Processor has to spend much of the time waiting while instructions and data are being fetched from main memory.
 - Memory speed cannot be increased beyond a certain point.



Impact of Processor / Memory Performance Gap

Year	CPU Clock	Clock Cycle	Memory Access	Minimum CPU Stall Cycles
1986	8 MHz	125 ns	190 ns	$190 / 125 - 1 = 0.5$
1989	33 MHz	30 ns	165 ns	$165 / 30 - 1 = 4.5$
1992	60 MHz	16.6 ns	120 ns	$120 / 16.6 - 1 = 6.2$
1996	200 MHz	5 ns	110 ns	$110 / 5 - 1 = 21.0$
1998	300 MHz	3.33 ns	100 ns	$100 / 3.33 - 1 = 29.0$
2000	1 GHz	1 ns	90 ns	$90 / 1 - 1 = 89.0$
2002	2 GHz	0.5 ns	80 ns	$80 / 0.5 - 1 = 159.0$
2004	3 GHz	0.33 ns	60 ns	$60 / 0.33 - 1 = 179.0$

Ideal memory access time
= 1 CPU cycle

Real memory access time
>> 1 CPU cycles

- **Memory Latency Reduction Techniques:**

- Faster DRAM cells (depends on VLSI technology)
- Wider memory bus width (fewer memory accesses needed)
- Multiple memory banks
- Integration of memory controller with processor
- New emerging RAM technologies

- **Memory Latency Hiding Techniques**

- Memory hierarchy (using SRAM-based cache memories)
- Pre-fetching instructions and/or data from memory before they are actually needed (used to hide long memory access latency)

Locality of Reference

- Programs tend to reuse data and instructions they have used recently.
 - **Rule of thumb:** 90% of the total execution time of a program is spent in only 10% of the code (also called 90/10 rule).
 - **Reason:** nested loops in a program, few procedures calling each other repeatedly, arrays of data items being accessed sequentially, etc.
- Basic idea to exploit this rule:
 - Based on a program's recent past, we can predict with a reasonable accuracy what instructions and data will be accessed in the near future.

- The 90/10 rule has two dimensions:
 - a) **Temporal Locality** (locality in time)
 - If an item is referenced in memory, it will tend to be referenced again soon.
 - b) **Spatial locality** (locality in space)
 - If an item is referenced in memory, nearby items will tend to be referenced soon.

(a) Temporal Locality

- Recently executed instructions are likely to be executed again very soon.
- Example: computing factorial of a number.

```
fact = 1;  
for k = 1 to N  
    fact = fact * k;
```



```
ADDI $t1,$zero,1  
ADDI $t2,$zero,N  
ADDI $t3,$zero,1  
Loop: MUL $t1,$t1,$t3  
ADDI $t3,$t3,1  
SGT $t4,$t3,$t2  
BNEZ $t4,Loop
```

- The four instructions in the loop are executed more frequently than the others.

(b) Spatial Locality

- Instructions residing close to a recently executing instruction are likely to be executed soon.
- Example: accessing elements of an array.

```
sum = 0;  
for k = 1 to N  
    sum = sum + A[k];
```



```
        SUB    $t1,$t1,$t1  
        ADDI   $t2,$zero,N  
        ADDI   $t3,$zero,1  
        ADDI   $t5,$zero,A  
Loop: LW     $t8,0($t5)  
        ADD    $t1,$t1,$t8  
        ADDI   $t3,$t3,1  
        SGT    $t4,$t3,$t2  
        BNEZ   $t4,Loop
```

- Performance can be improved by copying the array into cache memory.


- An example:
 - Accessing a 2-D array row-wise or column-wise.
 - Assume that the array elements are stored row-wise in memory.
 - One row of data can be brought into cache memory at any given time.

```
for (i=0; i<128; i++)
    for (j=0; j<128; j++)
        A[i][j] = A[i][j] + 1;
```

```
for (j=0; j<128; j++)
    for (i=0; i<128; i++)
        A[i][j] = A[i][j] + 1;
```

Performance of Memory Hierarchy

- We first consider a 2-level hierarchy consisting of two levels of memory, say, M_1 and M_2 .
 - M_1 can be cache memory, and M_2 can be main memory.



a) Cost:

- Let c_i denote the cost per bit of memory M_i , and S_i denote the storage capacity in bits of M_i .
- The average cost per bit of the memory hierarchy is given by:

$$\text{Cost } c = \frac{c_1 S_1 + c_2 S_2}{S_1 + S_2}$$

- In order to have $c \rightarrow c_2$, we must ensure that $S_1 \ll S_2$.

b) Hit Ratio / Hit Rate:

- The hit ratio H is defined as the probability that a logical address generated by the CPU refers to information stored in M_1 .
- We can determine H experimentally as follows:
 - A set of representative programs is executed or simulated.
 - The number of references to M_1 and M_2 , denoted by N_1 and N_2 respectively, are recorded.

$$H = \frac{N_1}{N_1 + N_2}$$

- The quantity $(1 - H)$ is called the *miss ratio*.

c) Access Time:

- Let t_{A1} and t_{A2} denote the access times of M_1 and M_2 respectively, relative to the CPU.
- The average time required by the CPU to access a word in memory can be expressed as:

$$t_A = H \cdot t_{A1} + (1 - H) \cdot t_{MISS}$$

where t_{MISS} denotes the time required to handle the miss, called *miss penalty*.

- The miss penalty t_{MISS} can be estimated in various ways:
 - The simplest approach is to set $t_{MISS} = t_{A2}$, that is, when there is a miss the data is accessed directly from M_2 .
 - A request for a word not in M_1 typically causes a block containing the requested word to be transferred from M_2 to M_1 . After completion of the block transfer, the word can be accessed in M_1 .

If t_B denotes the block transfer time, we can write

$$t_{MISS} = t_B + t_{A1} \quad [\text{since } t_B \gg t_{A1}, t_{A2} \approx t_B]$$

$$\text{Thus, } t_A = H \cdot t_{A1} + (1 - H) \cdot (t_B + t_{A1})$$

- c) If t_{HIT} denotes the time required to check whether there is a hit, we can write

$$t_{MISS} = t_{HIT} + t_B + t_{A1}$$

d) Efficiency:

- Let $r = t_{A2} / t_{A1}$ denote the access time ratio of the two levels of memory.
- We define the *access efficiency* as $e = t_{A1} / t_A$, which is the factor by which t_A differs from its minimum possible value.

$$\text{Efficiency } e = \frac{t_{A1}}{H \cdot t_{A1} + (1 - H) \cdot t_{A2}} = \frac{1}{H + (1 - H) \cdot r}$$

e) Speedup:

- The *speedup* gained by using the memory hierarchy is defined as $S = t_{A2} / t_A$.
- We can write:

$$S = \frac{t_{A2}}{H \cdot t_{A1} + (1 - H) \cdot t_{A2}} = \frac{1}{(1 - H) + H / r}$$

- The same result follows from *Amdahl's law*.

Some Common Terminologies Used

- **Block**: The smallest unit of information transferred between two levels.
- **Hit Rate**: The fraction of memory accesses found in the upper level.
- **Hit Time**: Time to access the upper level
 - Upper level access time + Time to determine hit/miss
- **Miss**: Data item needs to be retrieved from a block in the lower level.
- **Miss Rate**: The fraction of memory accesses not found in the upper level.
- **Miss Penalty**: Overhead whenever a miss occurs.
 - Time to replace a block in the upper level + Time to transfer the missed block

Example 1

- Consider a 2-level memory hierarchy consisting of a cache memory M_1 and the main memory M_2 . Suppose that the cache is 6 times faster than the main memory, and the cache can be used 90% of the time. How much speedup do we gain by using the cache?

Here, $r = 6$ and $H = 0.90$


$$\text{Thus, } S = 1 / [H / r + (1 - H)] = 1 / (0.90 / 6 + 0.10) = 1 / 0.25 = 4$$

Example 2


- Consider a 2-level memory hierarchy with separate instruction and data caches in level 1, and main memory in level 2.

The following parameters are given:

- The clock cycle time is 2 ns.
- The miss penalty is 15 clock cycles (for both read and write).
- 1 % of instructions are not found in I-cache.
- 8 % of data references are not found in D-cache.
- 20 % of the total memory accesses are for data.
- Cache access time (including hit detection) is 1 clock cycle.



Determine the overall average access time.



$$t_{MISS} = 1 + 15 = 16 \text{ cycles}$$


Average number of cycles per access:


$$\begin{aligned}
 & 0.80 \times (0.99 \times 1 + 0.01 \times 16) + 0.20 \times (0.92 \times 1 + 0.08 \times 16) \\
 & = 0.92 + 0.44 = 1.36
 \end{aligned}$$

Thus, average access time $t_A = 1.36 \times 2 \text{ ns} = 2.72 \text{ ns}$

Performance Calculation for Multi-Level Hierarchy

- Most of the practical memory systems use more than 2 levels of hierarchy.





t_{L1} : access time of M_1
 t_{L2} : access time of M_2
 H_{L1} : hit ratio of M_1
 H_{L2} : hit ratio of M_2 with respect to
 the residual accesses that try to
 access M_2

- Consider a 3-level hierarchy consisting of L1-cache, L2-cache and main memory.
 - Whenever there is a miss in $L1$, we go to $L2$.
 - Average access time can be calculated as:
- $$t_A = H_{L1} \cdot t_{L1} + (1 - H_{L1}) \cdot [H_{L2} \cdot t_{L2} + (1 - H_{L2}) \cdot t_{MISS2}]$$
- Here, t_{MISS2} is the miss penalty when the requested data is found neither in M_1 nor in M_2 .

Implications of a Memory Hierarchy to the CPU


- Processors designed without memory hierarchy are simpler because all memory accesses take the same amount of time.
 - Misses in a memory hierarchy implies *variable memory access times* as seen by the CPU.
- Some mechanism is required to determine whether or not the requested information is present in the top level of the memory hierarchy.
 - Check happens on every memory access and affects hit time.
 - Implemented in *hardware* to provide acceptable performance.

- Some mechanism is required to transfer blocks between consecutive levels.
 - If the block transfer requires 10's of clock cycles (like in cache / main memory hierarchy), it is controlled by *hardware*.
 - If the block transfer requires 1000's of clock cycles (like in main memory / secondary memory hierarchy), it can be controlled by *software*.
- Four main questions:
 1. *Block Placement*: Where to place a block in the upper level?
 2. *Block Identification*: How is a block found if present in the upper level?
 3. *Block Replacement*: Which block is to be replaced on a miss?
 4. *Write Strategy*: What happens on a write?

Cache Memory

Introduction

- Let us consider a single-level cache, and that part of the memory hierarchy consisting of cache memory and main memory.



- Cache memory is logically divided into *blocks* or *lines*, where every block (line) typically contains 8 to 256 bytes.
- When the CPU wants to access a word in memory, a special hardware first checks whether it is present in cache memory.
 - If so (called *cache hit*), the word is directly accessed from the cache memory.
 - If not, the block containing the requested word is brought from main memory to cache.
 - For writes, sometimes the CPU can also directly write to main memory.
- Objective is to keep the commonly used blocks in the cache memory.
 - Will result in significantly improved performance due to the property of *locality of reference*.

Q1. Where can a block be placed in the cache?

- This is determined by some *mapping algorithms*.
 - Specifies which main memory blocks can reside in which cache memory blocks.
 - At any given time, only a small subset of the main memory blocks can be held in the cache.
- Three common block mapping techniques are used:
 - a) **Direct Mapping**
 - b) **Associative Mapping**
 - c) **(N-way) Set Associative Mapping**

An Example Scenario: A 2-level memory hierarchy

- Consider a 2-level cache memory / main memory hierarchy.
 - The cache memory consists of 256 blocks (lines) of 32 words each.
 - Total cache size is 8192 (8K) words.
 - Main memory is addressable by a 24-bit address.
 - Main memory is word addressable.
 - Total size of the main memory is $2^{24} = 16 \text{ M}$ words.
 - Number of 32-word blocks in main memory = $16 \text{ M} / 32 = 512\text{K}$

(a) Direct Mapping

- Each main memory block can be placed in only one block in the cache.
- The mapping function is:

Cache Block = (Main Memory Block) % (Number of cache blocks)


- For the example,

Cache Block = (Main Memory Block) % 256


- Some example mappings:

$0 \rightarrow 0, 1 \rightarrow 1, 255 \rightarrow 255, 256 \rightarrow 0, 257 \rightarrow 1, 512 \rightarrow 0, 513 \rightarrow 1, \text{etc.}$

Direct Mapping



Direct Mapped Cache



- Block replacement algorithm is trivial, as there is no choice.
- More than one MM block is mapped onto the same cache block.
 - May lead to contention even if the cache is not full.
 - New block will replace the old block.
 - May lead to poor performance if both the blocks are frequently used.
- The MM address is divided into three fields: **TAG**, **BLOCK** and **WORD**.
 - When a new block is loaded into the cache, the 8-bit **BLOCK** field determines the cache block where it is to be stored.
 - The high-order 11 bits are stored in a **TAG** register associated with the cache block.
 - When accessing a memory word, the corresponding **TAG** field is compared.
 - Match implies **HIT**.


(b) Associative Mapping

- Here, a MM block can potentially reside in *any cache block position*.
- The memory address is divided into two fields: **TAG** and **WORD**.
 - When a block is loaded into the cache from MM, the higher order 19 bits of the address are stored into the **TAG** register corresponding to the cache block.
 - When accessing memory, the 19-bit **TAG** field of the address is compared with all the **TAG** registers corresponding to all the cache blocks.
- Requires *associative memory* for storing the **TAG** values.
 - High cost / lack of scalability.
- Because of complete freedom in block positioning, a wide range of replacement algorithms is possible.


What is an Associative Memory?

- It is a memory where data is accessed *by contents* rather than address.
 - There is circuitry to compare the applied data value with all the stored values in memory in parallel.
 - Wherever there is a match, the memory system will be returning the information.
 - Very expensive to build – only small capacity units are feasible.

Fully Associative Mapping





Fully Associative Cache Design



(c) N-way Set Associative Mapping


- A group of N consecutive blocks in the cache is called a *set*.
- This algorithm is a balance of *direct mapping* and *associative mapping*.
 - Like direct mapping, a MM block is mapped to a set.
 $\text{Set Number} = (\text{MM Block Number}) \% (\text{Number of Sets in Cache})$
 - The block can be placed anywhere within the set (there are N choices)
- The value of N is a design parameter:
 - $N = 1$:: same as *direct mapping*.
 - $N = \text{number of cache blocks}$:: same as *associative mapping*.
 - Typical values of N used in practice are: 2, 4 or 8.

4-way Set Associative Mapping



- Illustration for $N = 4$:
 - Number of sets in cache memory = 64.
 - Memory blocks are mapped to a set using *modulo-64 operation*.
 - Example: MM blocks 0, 64, 128, etc. all map to set 0, where they can occupy any of the four available positions.
- MM address is divided into three fields: TAG , SET and $WORD$.
 - The TAG field of the address must be associatively compared to the TAG fields of all the 4 blocks of the selected set.
 - This instead of requiring a single large associative memory, we need a number of very small associative memories only one of which will be used at a time.

K-way set-associative cache



Example 3

A 64 KB four-way set-associative cache is byte-addressable and contains 32 B lines. Memory addresses are 32 b wide.

- a) How wide are the tags in this cache?
- b) Which main memory addresses are mapped to set number 5?

Solution:

The number of sets in the cache = $64\text{KB}/(4 \times 32\text{B}) = 512$.

- a. Address (32 b) = 5 b byte offset + 9 b set index + 18 b tag
- b. Addresses that have their 9-bit set index equal to 5. These are of the general form $2^{14}a + 2^5 \times 5 + b$; e.g., 160-191, 16 554-16 575, ...

32-bit address	Tag	Set index	Offset
	18 bits	9 bits	5 bits
Tag width = $32 - 5 - 9 = 18$	Set size = $4 \times 32\text{B} = 128\text{B}$ Number of sets = $2^{16}/2^7 = 2^9$		Line width = $32\text{B} = 2^5\text{B}$

Q2. How is a block found if present in cache?

- Caches include a **TAG** associated with each cache block.
 - The TAG of every cache block where the block being requested may be present needs to be compared with the TAG field of the MM address.
 - All the possible tags are compared in parallel, as speed is important.
- Mapping Algorithms?
 - Direct mapping requires a *single comparison*.
 - Associative mapping requires a *full associative search* over *all the TAGs* corresponding to all cache blocks.
 - Set associative mapping requires a *limited associative search* over the *TAGs of only the selected set*.

- Use of *valid bit*:
 - There must be a way to know whether a cache block contains *valid* or *garbage* information.
 - A *valid bit* can be added to the *TAG*, which indicates whether the block contains valid data.
 - If the *valid bit is not set*, there is *no need to match* the corresponding *TAG*.

Q3. Which block should be replaced on a cache miss?

- With fully associative or set associative mapping, there can be several blocks to choose from for replacement when a miss occurs.
- Two primary strategies are used:
 - a) *Random*: The candidate block is selected randomly for replacement. This simple strategy tends to spread allocation uniformly.
 - b) *Least Recently Used (LRU)*: The block replaced is the one that has not been used for the longest period of time.
 - Makes use of a *corollary* of temporal locality:

If recently used blocks are likely to be used again, then the best candidate for replacement is the least recently used block

- To implement the LRU algorithm, the cache controller must track the LRU block as the computation proceeds.
- Example: Consider a 4-way set associative cache.
 - For tracking the LRU block within a set, we use a 2-bit counter with every block.
 - *When hit occurs:*
 - Counter of the referenced block is reset to 0.
 - Counters with values originally lower than the referenced one are incremented by 1, and all others remain unchanged.
 - *When miss occurs:*
 - If the set is not full, the counter associated with the new block loaded is set to 0, and all other counters are incremented by 1.
 - If the set is full, the block with counter value 3 is removed, the new block put in its place, and the counter set to 0. The other three counters are incremented by 1.

- It may be verified that the counter values of occupied blocks are all distinct.
- An example:

x	Block 0
x	Block 1
x	Block 2
x	Block 3

Initial

x	Block 0
x	Block 1
0	Block 2
x	Block 3

Miss: Block 2

0	Block 0
x	Block 1
1	Block 2
x	Block 3

Miss: Block 0

1	Block 0
x	Block 1
2	Block 2
0	Block 3

Miss: Block 3

2	Block 0
0	Block 1
3	Block 2
1	Block 3

Miss: Block 1

0	Block 0
1	Block 1
3	Block 2
2	Block 3

Hit: Block 0

1	Block 0
2	Block 1
0	Block 2
3	Block 3

Miss: Block 2

2	Block 0
3	Block 1
1	Block 2
0	Block 3

Hit: Block 3

2	Block 0
3	Block 1
0	Block 2
1	Block 3

Hit: Block 2

0	Block 0
3	Block 1
1	Block 2
2	Block 3

Hit: Block 0

1	Block 0
0	Block 1
2	Block 2
3	Block 3

Miss: Block 1

1	Block 0
0	Block 1
2	Block 2
3	Block 3

Hit: Block 1

Types of Cache Misses

a) Compulsory Miss

- On the first access to a block, the block must be brought into the cache.
- Also known as cold start misses, or first reference misses.
- Can be reduced by increasing cache block size or prefetching cache blocks.

b) Capacity Miss

- Blocks may be replaced from cache because the cache cannot hold all the blocks needed by a program.
- Can be reduced by increasing the total cache size.

c) Conflict Miss

- In case of direct mapping or N -way set associative mapping, several blocks may be mapped to the same block or set in the cache.
- May result in block replacements and hence access misses, even though all the cache blocks may not be occupied.
- Can be reduced by increasing the value of N (cache associativity).


Q4. What happens on a write?

- Statistical data suggests that read operations (including instruction fetches) dominate processor cache accesses.
 - All instruction fetch operations are read.
 - Most instructions do not write to memory.
- Making the *common case fast*:
 - Optimize cache accesses for reads.
 - But Amdahl's law reminds that for high performance designs we cannot ignore the speed of write operations.

- The common case (read operations) is relatively easy to make faster.
 - A block(s) can be read at the same time while the **TAG** is being compared with the block address.
 - If the read is a **HIT** the data can be passed to the CPU; if it is a **MISS** ignore it.
- Problems with write operations:
 - The CPU specifies the size of the write (between 1 and 8 bytes), and only that portion of a block has to be changed.
 - Implies a **read-modify-write** sequence of operations on the block.
 - Also, the process of modifying the block cannot begin until the TAG is checked to see if it is a hit.
 - Thus, cache write operations take more time than cache read operations.


Cache Write Strategies

- Cache designs can be classified based on the write and memory update strategy being used.
 1. *Write Through / Store Through*
 2. *Write Back / Copy Back*



(a) Write Through Strategy

- Information is written to both the cache block and the main memory block.
- Features:
 - Easier to implement.
 - Read misses do not result in writes to the lower level (i.e. MM).
 - The lower level (i.e. MM) has the most updated version of the data – important for I/O operations and multiprocessor systems.
 - A write buffer is often used to reduce CPU write stall time while data is written to main memory.



- **Perfect Write Buffer:**

- All writes are handled by write buffer; no stalling for write operations.
- For unified L1 cache,

$$\text{Stall Cycles / Memory Access} = \% \text{ Reads} \times (1 - H_{L1}) \cdot t_{MM}$$


- **Realistic Write Buffer:**

- A percentage of write stalls are not eliminated when the write buffer is full.
- For unified L1 cache,

$$\text{Stall Cycles / Memory Access} = (\% \text{ Reads} \times (1 - H_{L1}) + \% \text{ write stalls not eliminated}) \times t_{MM}$$

(b) Write Back Strategy

- Information is written only to the cache block.
- A modified cache block is written to MM only when it is replaced.
- Features:
 - Writes occur at the speed of cache memory.
 - Multiple writes to a cache block requires only one write to MM.
 - Uses less memory bandwidth, makes it attractive to multiprocessors.
- Write-back cache blocks can be *clean* or *dirty*.
 - A status bit called *dirty bit* or *modified bit* is associated with each cache block, which indicates whether the block was modified in the cache (0: clean, 1: dirty).
 - If the status is clean, the block is not written back to MM while being replaced.



Cache Write Miss Policy

- Since information is usually not needed immediately on a write miss, two options are possible on a cache write miss:

a) Write Allocate

- The missed block is loaded into cache on a write miss, followed by write hit actions.
- Requires a cache block to be *allocated* for the block to be written into.

b) No-Write Allocate

- The block is modified only in the lower level (i.e. MM), and not loaded into cache.
- Cache block is *not allocated* for the block to be written into.

- Typical usage:
 - a) Write-back cache with write-allocate
 - In order to capture subsequent writes to the block in cache.
 - b) Write-through cache with no-write-allocate
 - Since subsequent writes still have to go to MM.

Estimation of Miss Penalties

- Write-Through Cache
 - Write Hit Operation:
 - Without write buffer, miss penalty = t_{MM}
 - With perfect write buffer, miss penalty = 0
- Write-Back Cache
 - Write Hit Operation
 - Miss penalty = 0

- **Write-Back Cache (with Write Allocate)**

- **Write Hit Operation**

- Miss penalty = O

- **Read or Write Miss Operation**

- If the replaced block is clean, miss penalty = t_{MM}

- No need to write the block back to MM.
 - New block to be brought into MM (t_{MM}).

- If the replaced block is dirty, miss penalty = $2 t_{MM}$

- Write the block to be replaced to MM (t_{MM}).
 - New block to be brought into MM (t_{MM}).

Example 4

- Consider a CPU with average CPI of 1.1.
 - Assume an instruction mix: ALU – 50%, LOAD – 15%, STORE – 15%, BRANCH – 20%
 - Assume a cache miss rate of 1.5%, and miss penalty of 50 cycles ($= t_{MM}$).
 - Calculate the effective CPI for a unified L1 cache, using *write through and no write allocate*, with:
 - a) No write buffer
 - b) Perfect write buffer
 - c) Realistic write buffer that eliminates 85% of write stalls.

Number of memory accesses per instruction = $1 + 15\% + 15\% = 1.3$

% Reads = $(1 + 0.15) / 1.3 = 88.5\%$ % Writes = $0.15 / 1.3 = 11.5\%$

- **Solution:**

- a) With no write buffer (i.e. stall on all writes)
 - Memory stalls / instr. = $1.3 \times 50 \times (88.5\% \times 1.5\% + 11.5\%) = 8.33$ cycles
 - CPI = $CPI_{avg} + \text{Memory stalls / instr.} = 1.1 + 8.33 = 9.43$
- b) With perfect write buffer (i.e. all write stalls are eliminated)
 - Memory stalls / instr. = $1.3 \times 50 \times (88.5\% \times 1.5\%) = 0.86$ cycles
 - CPI = $1.1 + 0.86 = 1.96$
- c) With realistic write buffer (85% of write stalls are eliminated)
 - Memory stalls / instr. = $1.3 \times 50 \times (88.5\% \times 1.5\% + 15\% \times 11.5\%) = 1.98$ cycles
 - CPI = $1.1 + 1.98 = 3.08$

Example 5

- Consider a CPU with average CPI of 1.1.
 - Assume the instruction mix: ALU – 50%, LOAD – 15%, STORE – 15%, BRANCH – 20%
 - Assume a cache miss rate of 1.5%, and miss penalty of 50 cycles ($= t_{MM}$).
 - Calculate the effective CPI for a unified L1 cache, using *write back and write allocate*, with the probability of a cache block being dirty is 10%.

Number of memory accesses per instruction = $1 + 15\% + 15\% = 1.3$

- **Solution:**

- Memory accesses per instruction = 1.3
- Stalls / access = $(1 - H_{L1}) \cdot (t_{MM} \times \% \text{ clean} + 2t_{MM} \times \% \text{ dirty})$
= $1.5\% \times (50 \times 90\% + 100 \times 10\%) = 0.825 \text{ cycles}$
- Average memory access time = $1 + \text{stalls / access} = 1 + 0.825 = 1.825 \text{ cycles}$
- Memory stalls / instr. = $1.3 \times 0.825 = 1.07 \text{ cycles}$
- Thus, effective CPI = $1.1 + 1.07 = 2.17$

Improving Cache Performance

Introduction


- We shall discuss various techniques using which the performance of cache memory can be improved.
- We consider the following expression for *average memory access time* (AMAT):
$$AMAT = Hit\ time + Miss\ rate \times Miss\ penalty$$
- When we talk about improving the performance of cache memory systems, we can try to reduce one or more of the three parameters: *Hit time, Miss rate, Miss penalty*.

Basic Cache Optimization Techniques

- We can categorize the techniques into three categories based on the parameter that is being optimized:
 - a) ***Reducing the miss rate***: we can use larger block size, larger cache size, and higher associativity.
 - b) ***Reducing the miss penalty***: we can use multi-level caches and giving priority to reads over writes.
 - c) ***Reducing the cache hit time***: we can avoid the address translation when indexing the cache.

(a) Use Larger Block Size

- Increasing the block size helps in reducing the miss rate.
 - See plot on the next slide.
- Larger blocks also reduce compulsory misses.
 - Since larger blocks can take better advantage of spatial locality.
- **Drawbacks:**
 - The miss penalty increases, as it is required to transfer larger blocks.
 - Since the number of cache blocks decreases, the number of conflict misses and even capacity misses can increase.
 - The overheads may outweigh the gain.




[Hennessy & Patterson, “Computer Architecture: A Quantitative Approach” (4/e)]

- **Selection of block size:**

- The optimal selection of the block size depends on both the latency and the bandwidth of the lower-level memory.
- **High latency and high bandwidth**
 - Encourages large block size since the cache gets many more bytes for a miss for a nominal increase in miss penalty.
- **Low latency and low bandwidth**
 - Encourages smaller block sizes since more time is required to transfer larger blocks.
 - Larger number of smaller blocks may also reduce conflict misses.

(b) Use Larger Cache Memory

- Increasing the size of the cache is a straightforward way to reduce the capacity misses.
- **Drawbacks:**
 - Increases the hit time since the number of TAGs to be searched in parallel will be possibly larger.
 - Results in higher cost and power consumption.
- Traditionally popular for off-chip caches.



[Hennessy & Patterson, “Computer Architecture: A Quantitative Approach” (4/e)]

(c) Use Higher Associativity

- For N -way associative cache, the miss rate reduces as we increase N .
 - Reduces conflict misses, as there are more choices to place a block in cache.
- General rule of thumb:
 - 8-way set associative cache is as effective as fully associative for practical scenarios.
 - A direct mapped cache of size N has about the same miss rate as a 2-way set associative cache of size $N/2$.
- **Drawbacks:**
 - Increases the hit time as we have to search a larger associative memory.
 - Increases power consumption due to higher complexity of associative memory.

(d) Use Multi-level Caches

- Here we try to reduce the miss penalty, and not the miss rate.
- Performance gap between processors and memory increases with time.
 - Use faster cache to keep pace with the speed of the processor.
 - Make the cache larger to bridge the widening gap between processor and MM.
- We can use both in a multi-level cache system:
 - The *L1* cache can be small enough to match the clock cycle time of the fast processor.
 - The *L2* cache can be large enough to capture many accesses that would go to MM, thereby reducing the miss penalty.

- Consider a 2-level cache system, consisting of L1 cache and L2 cache.
- The average memory access time can be computed as:

$$AMAT = HitTime_{L1} + MissRate_{L1} \times MissPenalty_{L1}$$

where $MissPenalty_{L1} = HitTime_{L2} + MissRate_{L2} \times MissPenalty_{L2}$

- Thus,

$$AMAT = HitTime_{L1} + MissRate_{L1} \times (HitTime_{L2} + MissRate_{L2} \times MissPenalty_{L2})$$

- The second-level miss rate $MissRate_{L2}$ is measured on the *leftovers* from the first-level cache.

- We define the following for a 2-level cache system:
 - **Local Miss Rate:**
 - This is defined as the number of misses in a cache divided by the total number of accesses to this cache.
 - For the first level, this is MissRate_{L1}
 - For the second level, this is MissRate_{L2}
 - **Global Miss Rate:**
 - This is defined as the number of misses in a cache divided by the total number of memory accesses generated by the processor.
 - For the first level, this is MissRate_{L1}
 - For the second level, this is $\text{MissRate}_{L1} \times \text{MissRate}_{L2}$

- The local miss rate is large for $L2$ cache because the $L1$ cache takes out a major fraction of the total memory accesses.
- For this purpose, the global miss rate is a more useful measure.
 - Fraction of memory accesses generated by the processor that goes all the way to main memory.
- A useful measure:

$$\begin{aligned} \text{Average Memory Stalls per Instr.} = & \text{ Misses-per-instr}_{L1} \times \text{HitTime}_{L2} \\ & + \text{ Misses-per-instr}_{L2} \times \text{MissPenalty}_{L2} \end{aligned}$$

Example 6

- Suppose that in 1000 memory references there are 60 misses in L1-cache and 15 misses in L2-cache. What are the various miss rates?
Assume that MissPenalty_{L2} is 180 clock cycles, HitTime_{L1} is 1 clock cycle, and HitTime_{L2} is 12 clock cycles.
What will be the average memory access time? Ignore the impact of writes.

Solution:

$$\text{MissRate}_{L1} = 60 / 1000 = 6\% \quad (\text{both local or global})$$

$$\text{LocalMissRate}_{L2} = 15 / 60 = 25\%$$

$$\text{GlobalMissRate}_{L2} = 15 / 1000 = 1.5\%$$

$$\begin{aligned}\text{AMAT} &= \text{HitTime}_{L1} + \text{MissRate}_{L1} \times (\text{HitTime}_{L2} + \text{MissRate}_{L2} \times \text{MissPenalty}_{L2}) \\ &= 1 + 6\% \times (12 + 25\% \times 180) \\ &= 1 + 6\% \times 57 = 4.42 \text{ clock cycles}\end{aligned}$$

- **Multi-level inclusion** versus **Multi-level exclusion**
 - **Multi-level inclusion** requires that $L1$ data are always present in $L2$.
 - Desirable because consistency between I/O and caches can be determined just by checking the $L2$ cache.
 - **Multi-level exclusion** requires that $L1$ data is *never* found in $L2$.
 - Typically, a cache miss in $L1$ results in a *swap* of blocks between $L1$ and $L2$ rather than a replacement of a $L1$ block with a $L2$ block.
 - This policy prevents wasting space in the $L2$ cache.
 - May make sense if the designer can only afford a $L2$ cache that is *slightly bigger* than the $L1$ cache.

(e) Giving Priority to Read Misses Over Writes

- The presence of *write buffers* can complicate memory accesses.
 - The buffer may be holding the updated value of a location needed on a read miss.
- Simplest solution is to make the read miss to wait until the write buffer is empty.
 - As an alternative, check the contents of the write buffer for any conflict; and if none, the read miss can continue → *reduces read miss penalty*.
 - Most desktops and servers follow this approach, giving priority to reads over writes.